


Dashboard Interpretation



The monitoring dashboard provides a comprehensive view of the operational performance of the closed-domain RAG system under simulated load. The request rate remains stable at approximately 0.2 requests per second, indicating a controlled and consistent traffic pattern. The steady increase in total requests confirms that Prometheus is reliably scraping and aggregating system metrics over time.

Latency metrics reveal moderate variability in system responsiveness. Average latency fluctuates between approximately 0.7 and 1.0 seconds, while P95 latency ranges from about 1.6 to 1.9 seconds. The gap between average and P95 latency suggests the presence of tail latency, where a subset of requests experiences slower processing times. This indicates potential inefficiencies in retrieval or generation steps that could impact user experience under higher load.

The retrieval top score varies between 0.7 and 0.95, demonstrating inconsistent retrieval quality. While high scores indicate strong semantic matches, lower values suggest occasional retrieval of less relevant documents. This variability directly affects downstream answer quality and grounding.

Groundedness scores generally remain above 0.8, indicating that most responses are supported by retrieved context. However, the grounded answer rate fluctuates significantly, ranging from approximately 40% to 80%. This inconsistency suggests that even when relevant documents are retrieved, the generation component does not always produce fully grounded responses, highlighting a gap between retrieval quality and answer synthesis.

The refusal rate ranges from 0% to 30%, reflecting the system's handling of out-of-scope queries. This variability is expected given the mixed input distribution but also indicates that refusal behavior is not fully calibrated. A more stable refusal policy would improve reliability and user trust.

The error rate remains near zero throughout the observation period, indicating that the system is stable and free from runtime failures. This suggests that the monitoring and execution pipeline is robust, even if performance characteristics vary.